

Part 1:

Build the best network you can using Keras Summarize Accuracy measures for Train, Test and Validation for the best model

Architecture:

Layer (type)	Output Shape	Param #
conv2d_19 (Conv2D)	(None, 32, 32, 16)	448
activation_25 (Activation)	(None, 32, 32, 16)	0
batch_normalization_19 (Batch Normalization)	(None, 32, 32, 16)	64
conv2d_20 (Conv2D)	(None, 30, 30, 16)	2320
activation_26 (Activation)	(None, 30, 30, 16)	0
batch_normalization_20 (Batch Normalization)	(None, 30, 30, 16)	64
max_pooling2d_10 (MaxPooling2D)	(None, 15, 15, 16)	0
dropout_13 (Dropout)	(None, 15, 15, 16)	0
conv2d_21 (Conv2D)	(None, 15, 15, 32)	4640
activation_27 (Activation)	(None, 15, 15, 32)	0
batch_normalization_21 (Batch Normalization)	(None, 15, 15, 32)	128
conv2d_22 (Conv2D)	(None, 13, 13, 32)	9248
activation_28 (Activation)	(None, 13, 13, 32)	0
batch_normalization_22 (Batch Normalization)	(None, 13, 13, 32)	128
max_pooling2d_11 (MaxPooling2D)	(None, 6, 6, 32)	0
dropout_14 (Dropout)	(None, 6, 6, 32)	0
conv2d_23 (Conv2D)	(None, 6, 6, 64)	18496

activation_29 (Activation)	(None, 6, 6, 64)	0
batch_normalization_23 (Batch Normalization)	(None, 6, 6, 64)	256
conv2d_24 (Conv2D)	(None, 4, 4, 64)	36928
activation_30 (Activation)	(None, 4, 4, 64)	0
batch_normalization_24 (Batch Normalization)	(None, 4, 4, 64)	256
max_pooling2d_12 (MaxPooling2D)	(None, 2, 2, 64)	0
dropout_15 (Dropout)	(None, 2, 2, 64)	0
conv2d_25 (Conv2D)	(None, 2, 2, 64)	36928
activation_31 (Activation)	(None, 2, 2, 64)	0
batch_normalization_25 (Batch Normalization)	(None, 2, 2, 64)	256
max_pooling2d_13 (MaxPooling2D)	(None, 1, 1, 64)	0
dropout_16 (Dropout)	(None, 1, 1, 64)	0
global_average_pooling2d_4 (GlobalAveragePooling2D)	(None, 64)	0
activation_32 (Activation)	(None, 64)	0
dropout_17 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 200)	13000
activation_33 (Activation)	(None, 200)	0


This is a general keras model that worked better than the others. We tried different parameter values and layer structure, though they actually did not make a pretty obvious difference.

We reformatted the dataset which can be delivered to the model, and set the validation split to 0.3.

Training Result:

Test loss: 3.8896696652730305



Test accuracy: 0.15346666666666667



Here is another example of our trainings:

Test loss: 3.8188165491739907

Test accuracy: 0.17563333333333334



Though the test accuracy is higher, but this model is too underfitting. So we changed the dropout and adjusted the layers counts and the result got a bit better.

Part 2:

Use Autokeras to tune hyperparameters. Summarize Accuracy measures for Train, Test and Validation for the best model.

Configuration:

We split the validation data as 30% of the whole dataset.

The training_times is 3 hours. The whole training progress spent around 12 hours.

Training Result:

Autokeras trained three models totally and gave us a big model containing 26 conv2D layers.

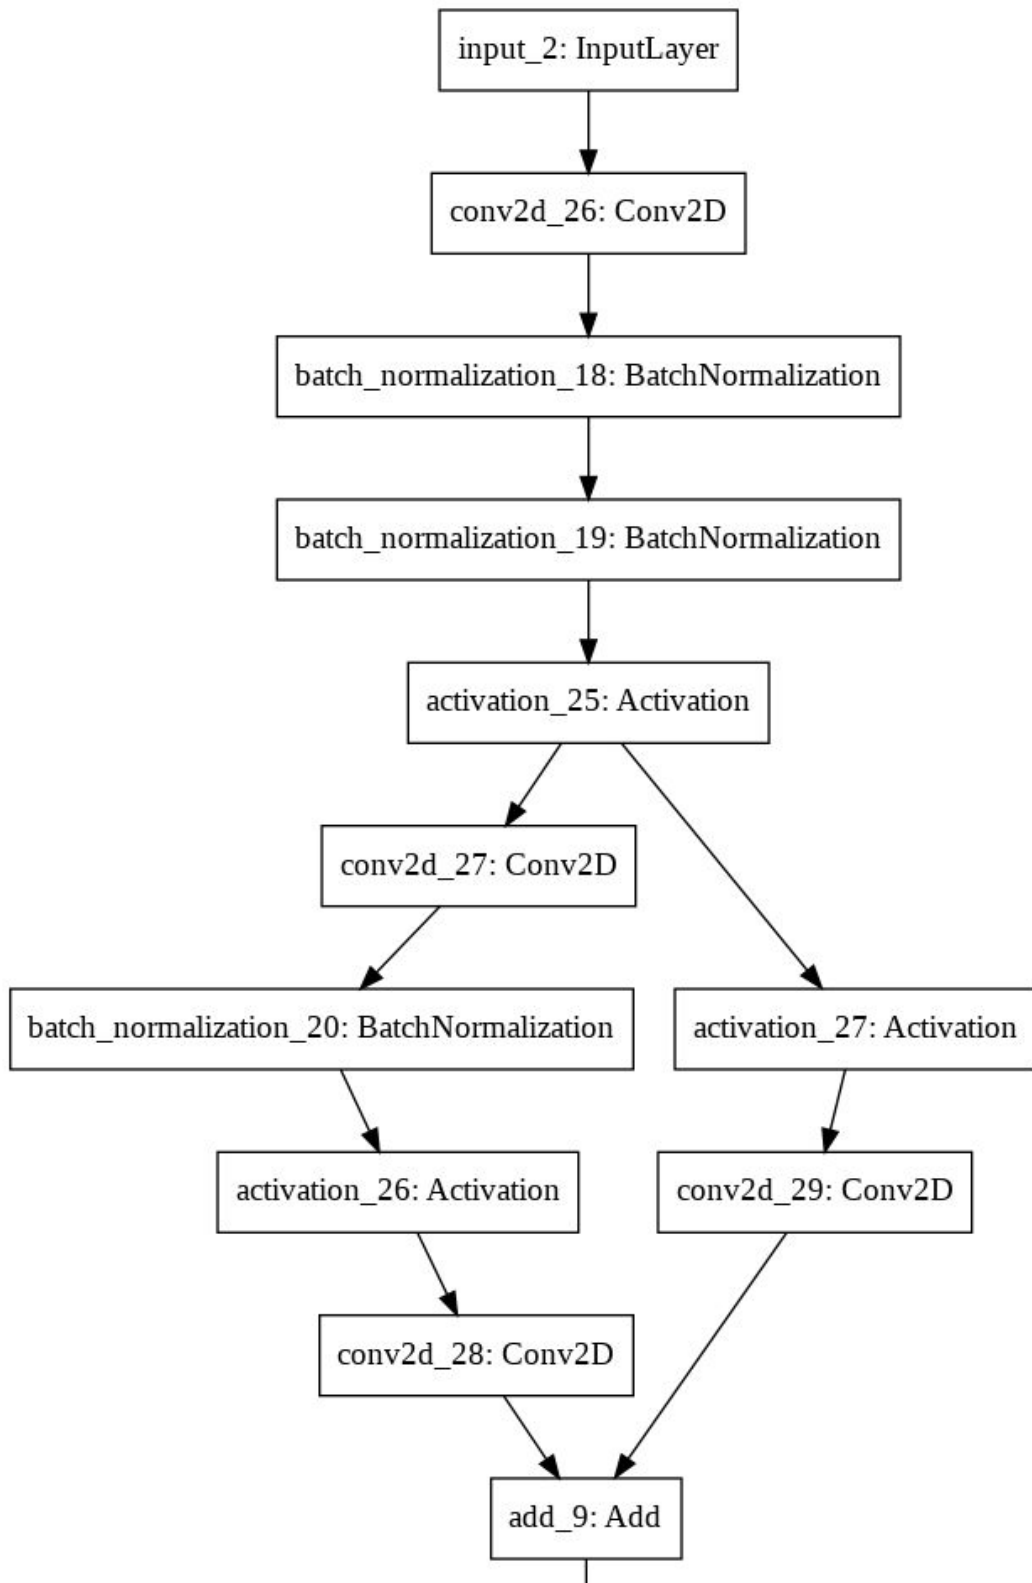
Here is the Autokeras report:

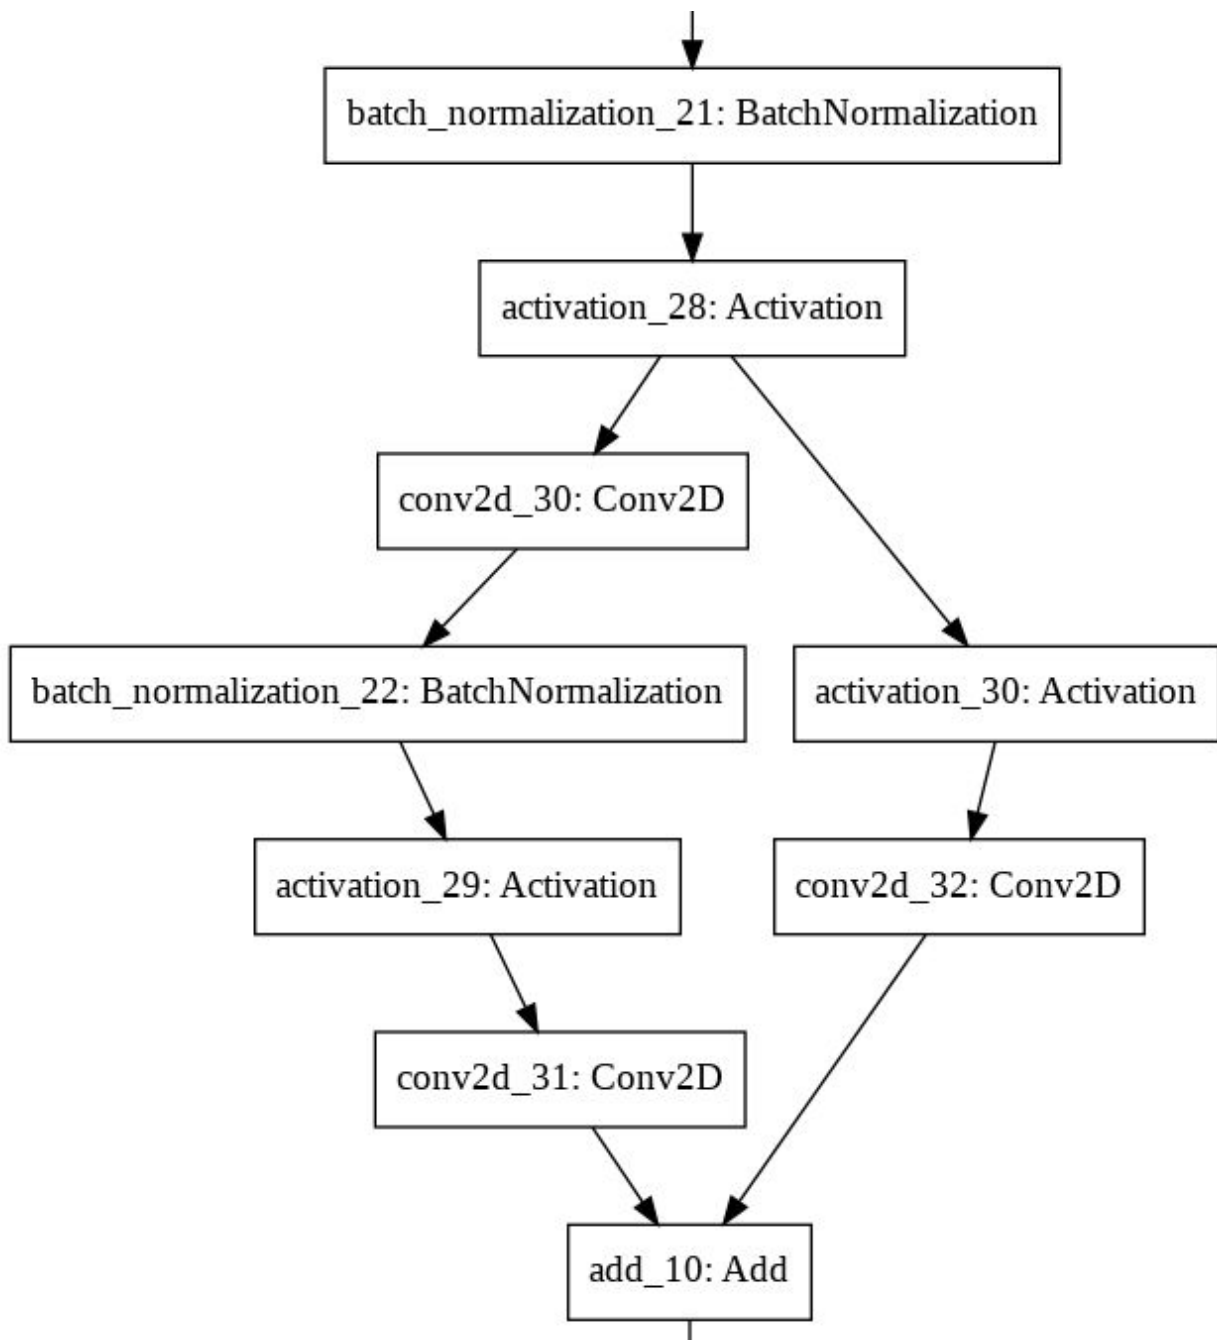
	precision	recall	f1-score	support
micro avg	0.30	0.30	0.30	30000
macro avg	0.34	0.30	0.29	30000
weighted avg	0.34	0.30	0.29	30000

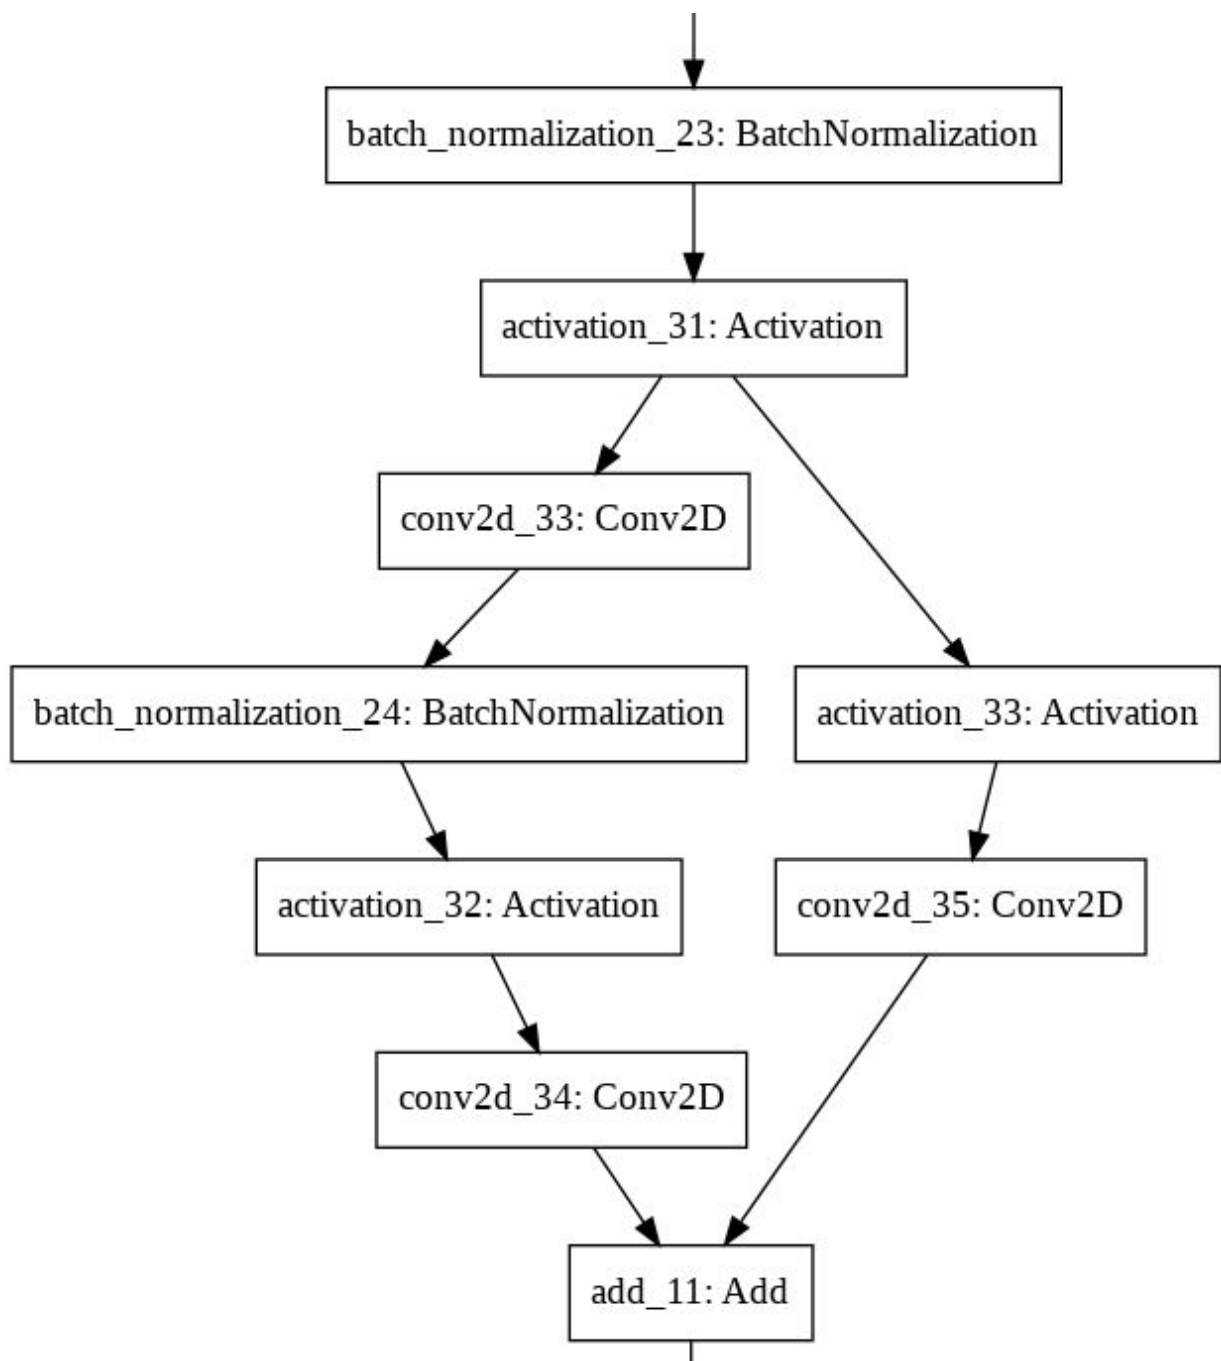
score: 0.3004666666666666

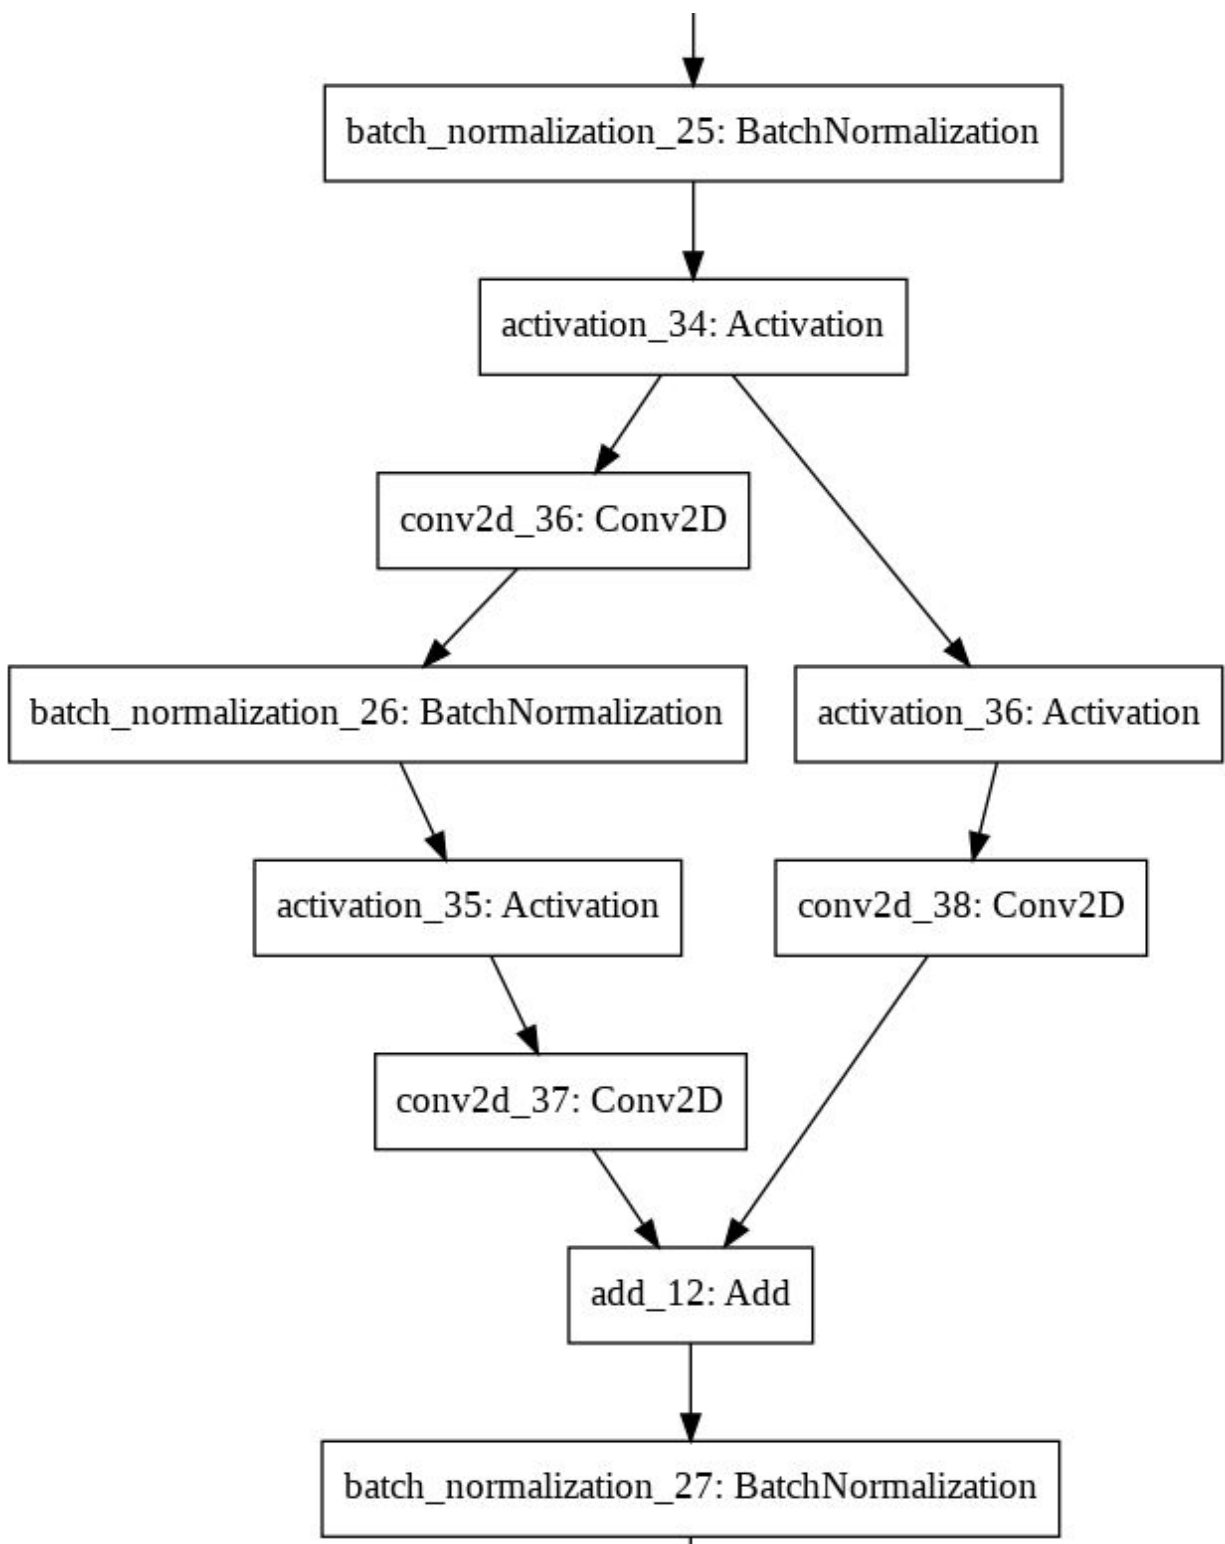
The final score is 0.3, and average accuracy is 0.34.

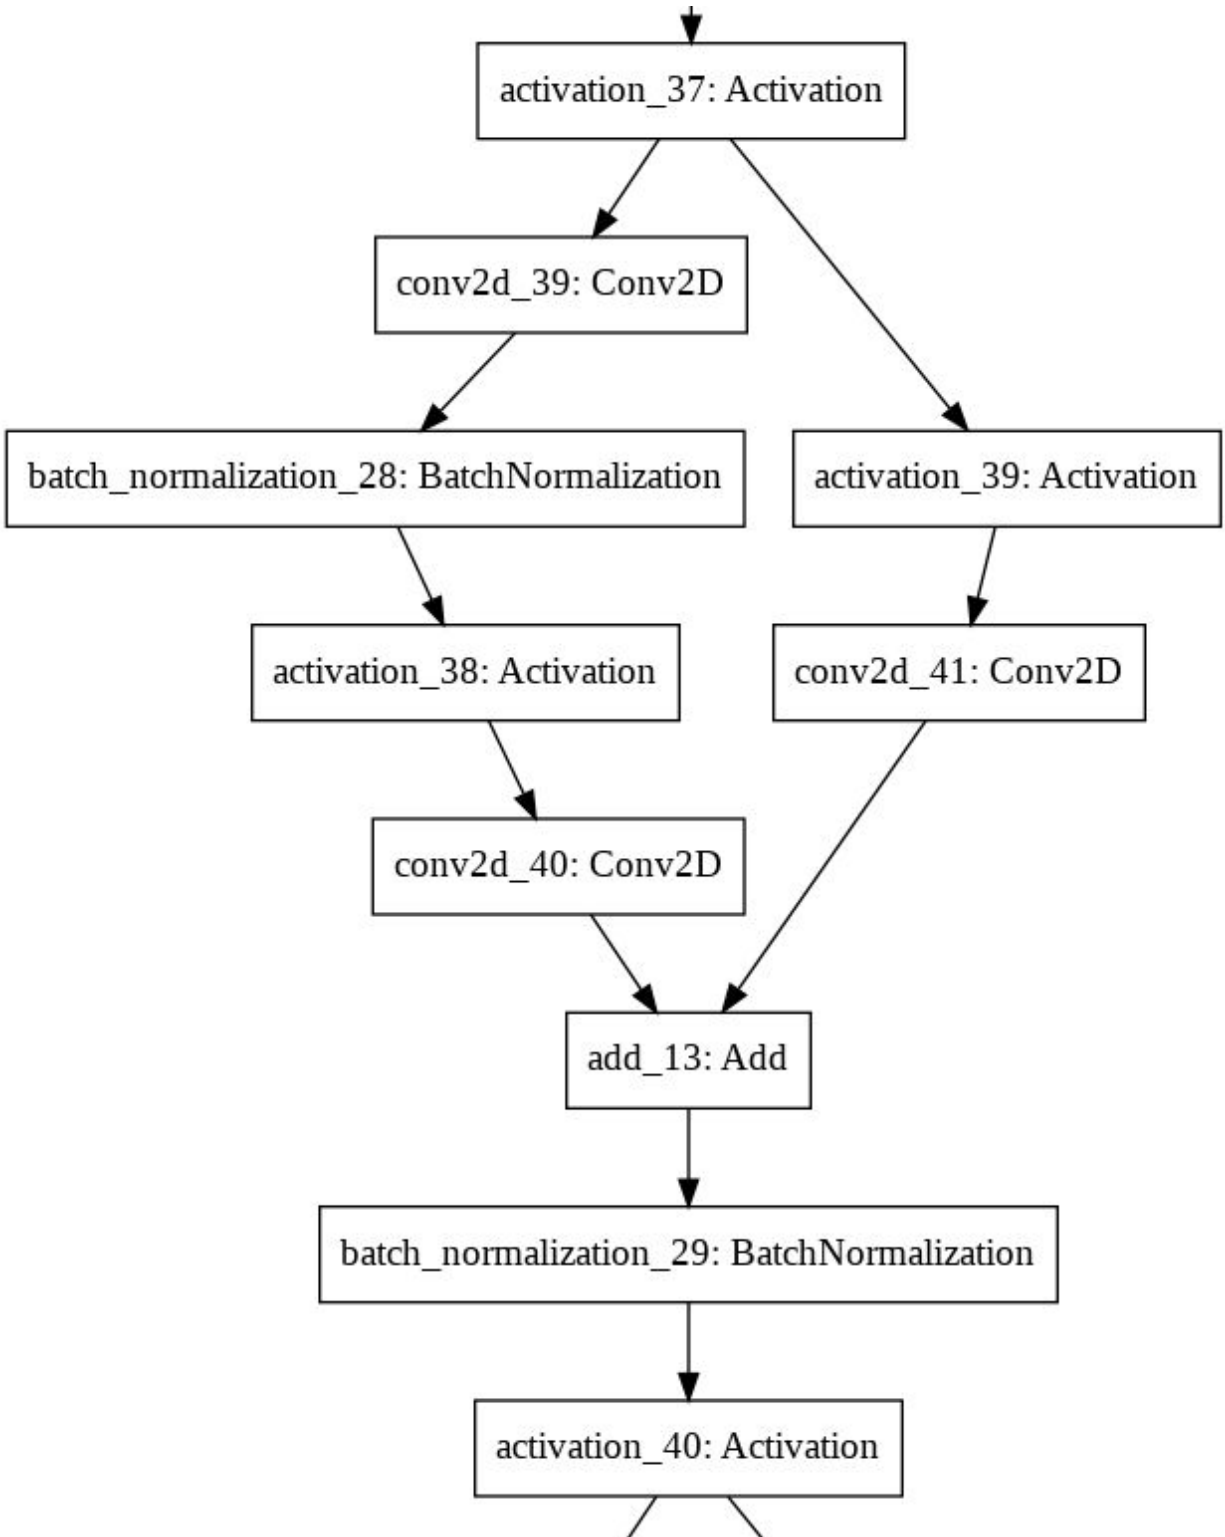
The following is the Autokeras model plot:

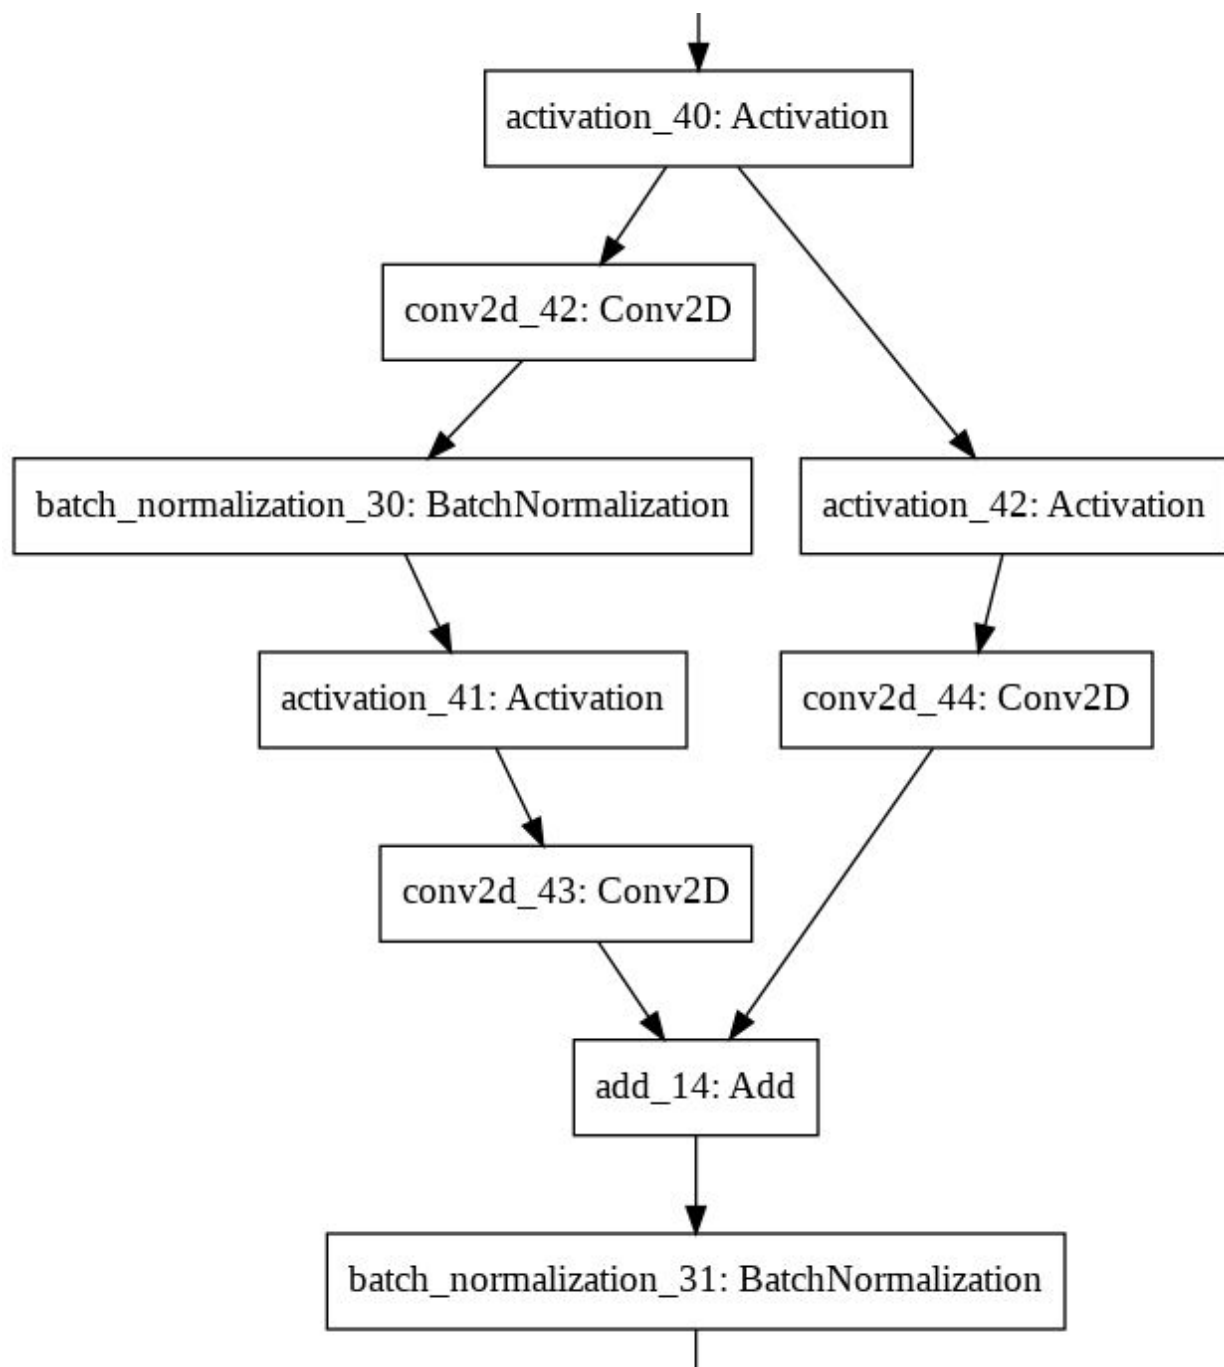


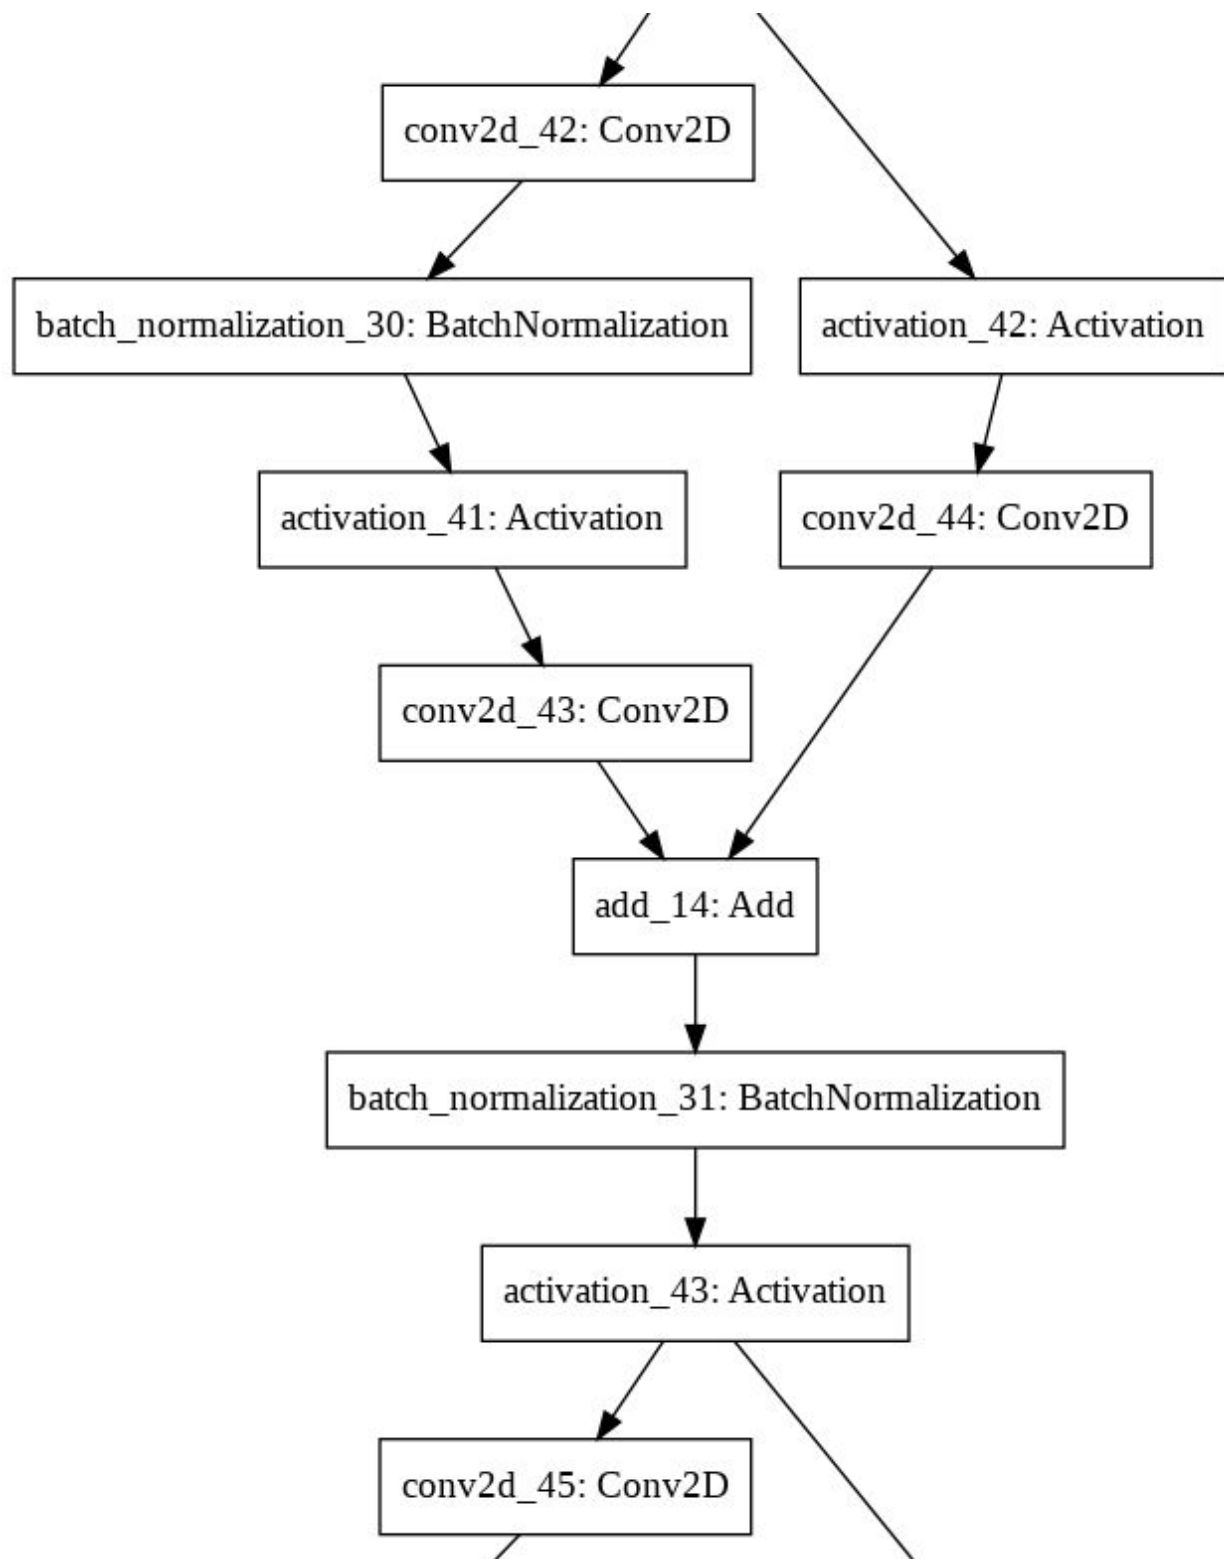


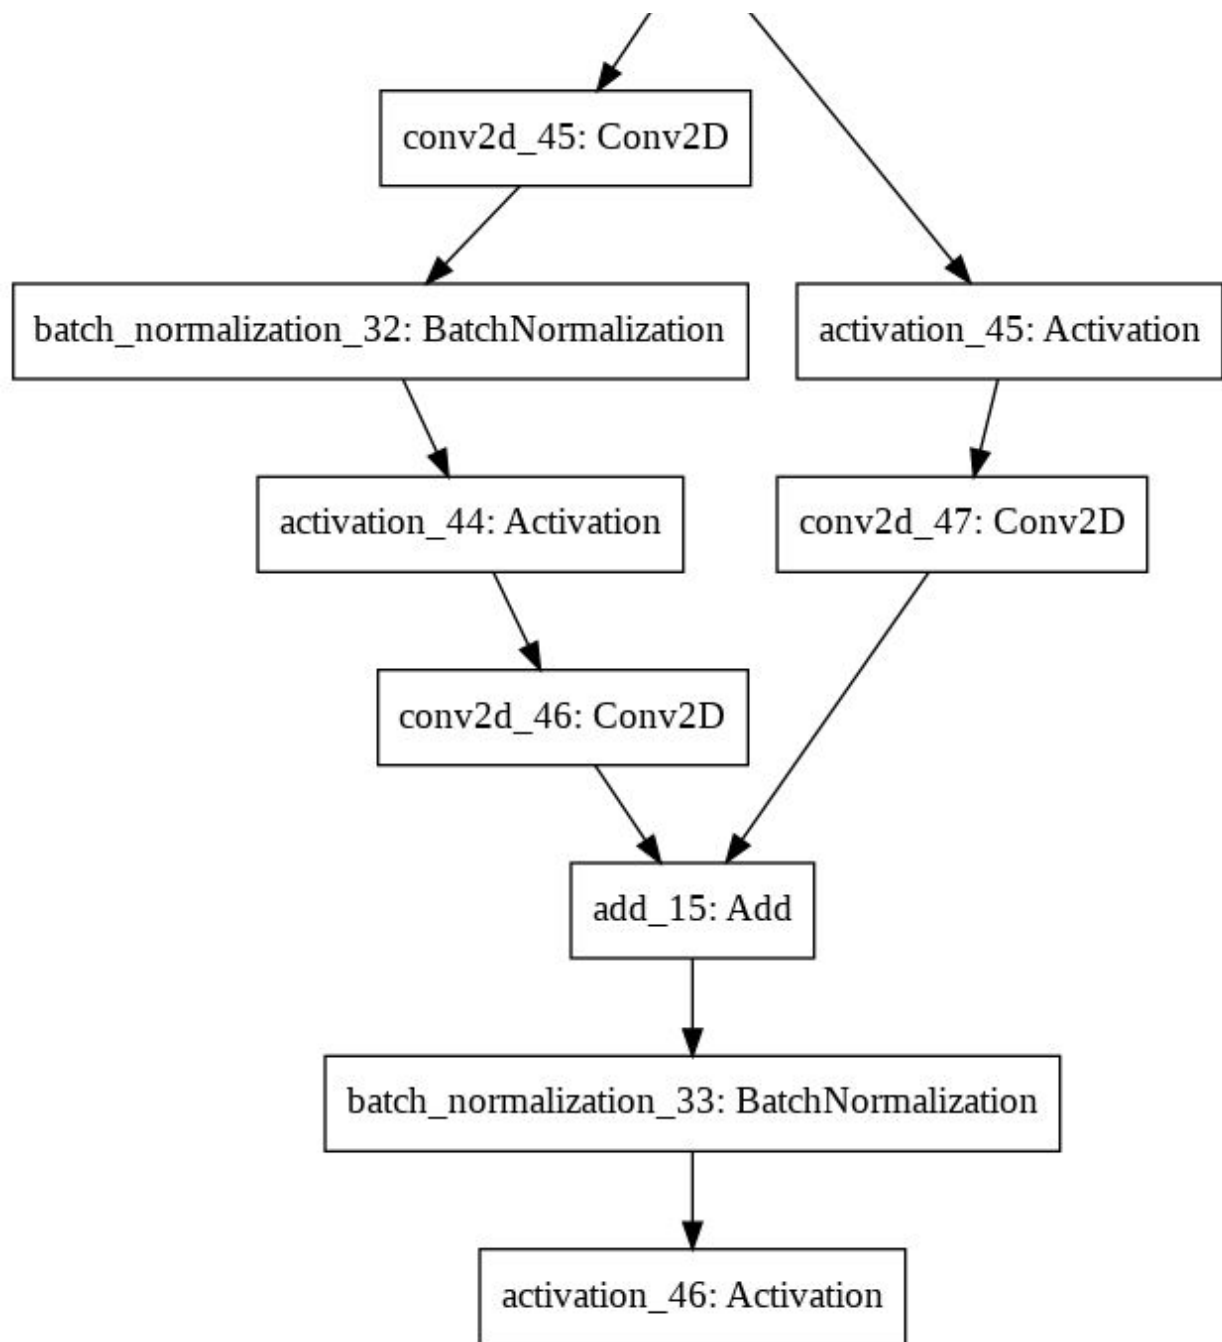


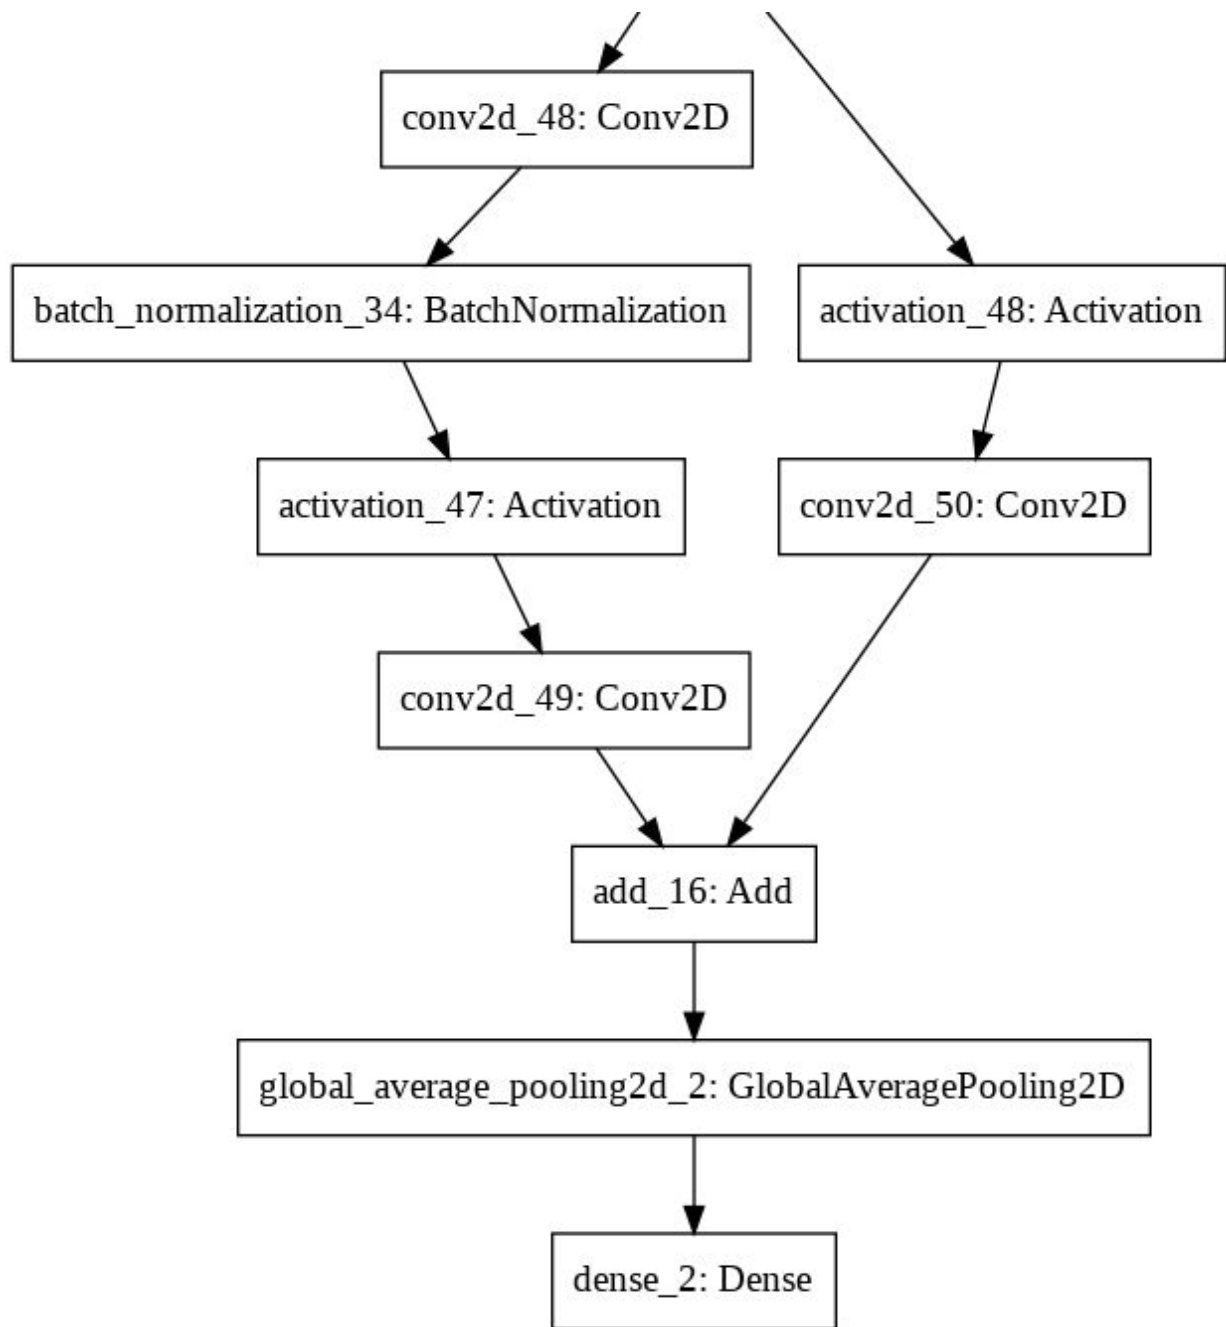












Implement the Alexnet models in Keras. Summarize Accuracy measures for Train, Test and Validation for the best model

```
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)
```


(70000, 32, 32, 3)
(70000, 200)
(30000, 32, 32, 3)
(30000, 200)

Data:

We implement the Alexnet model in Keras with some modification on our “Tiny-Imagenet/train” data set. The portion of train and test data set is 7 : 3 which are 70000 and 30000 pictures correspondingly.

Architecture:

The AlexNet architecture consists of five convolutional layers, some of which are followed by maximum pooling layers and then three fully-connected layers and finally a 1000-way softmax classifier.



In the original paper, all the layers are divided into two to train them on separate GPUs. Since it is a complex arrangement and difficult to understand, we will implement AlexNet model in one layer concept.

First Layer:

The input for AlexNet is a 227x227x3 RGB image which passes through the first convolutional layer with 96 feature maps or filters having size 11×11 and a stride of 4.

The image dimensions changes to 55x55x96.

$(227-11)/4+1=55$ (add one more pixel from the 11x11)

Then the AlexNet applies maximum pooling layer or sub-sampling layer with a filter size 3×3 and a stride of two. The resulting image dimensions will be reduced to 27x27x96.
 $(55-3)/2+1=27$

Our First Layer:

```
model.add(Conv2D(96, (3,3), strides=(2,2), activation='relu', padding='same', input_shape=(x_train.shape[1],x_train.shape[2],x_train.shape[3])))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2,2)))
# Local Response normalization for Original Alexnet
model.add(BatchNormalization())
```

conv2d_1 (Conv2D)	(None, 16, 16, 96)	2688
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 96)	0
batch_normalization_1 (Batch Normalization)	(None, 8, 8, 96)	384

Our input shape is 32x32x3 after reshape the “Tiny-Imagenet” image size, and with a stride of 2, the dimensions change to 16x16x96. Finally, we add one more batch_normalization layer after the pooling layer.

Second Layer:

Next, there is a second convolutional layer with 256 feature maps having size 5×5 and a stride of 1.

Then there is again a maximum pooling layer with filter size 3×3 and a stride of 2. This layer is same as the second layer except it has 256 feature maps so the output will be reduced to 13x13x256.

Our Second Layer:

```
model.add(Conv2D(256, (5,5), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(3, 3), strides=(2,2)))
# Local Response normalization for Original Alexnet
model.add(BatchNormalization())
```

conv2d_2 (Conv2D)	(None, 8, 8, 256)	614656
max_pooling2d_2 (MaxPooling2D)	(None, 3, 3, 256)	0
batch_normalization_2 (Batch Normalization)	(None, 3, 3, 256)	1024

Third, Fourth and Fifth Layers:

The third, fourth and fifth layers are convolutional layers with filter size 3×3 and a stride of one. The first two used 384 feature maps where the third used 256 filters.

The three convolutional layers are followed by a maximum pooling layer with filter size 3×3, a stride of 2 and have 256 feature maps.

Our Third, Fourth and Fifth Layers:

```
model.add(Conv2D(384, (3,3), activation='relu', padding='same'))
model.add(Conv2D(384, (3,3), activation='relu', padding='same'))
model.add(Conv2D(256, (3,3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(3, 3), strides=(2,2)))
# Local Response normalization for Original Alexnet
model.add(BatchNormalization())
```

conv2d_3 (Conv2D)	(None, 3, 3, 384)	885120
conv2d_4 (Conv2D)	(None, 3, 3, 384)	1327488
conv2d_5 (Conv2D)	(None, 3, 3, 256)	884992
max_pooling2d_3 (MaxPooling2	(None, 1, 1, 256)	0
batch_normalization_3 (Batch	(None, 1, 1, 256)	1024

Sixth Layer:

The convolutional layer output is flatten through a fully connected layer with 9216 feature maps each of size 1×1.

```
model.add(Flatten())
```

flatten_1 (Flatten)	(None, 256)	0
---------------------	-------------	---

Seventh and Eighth Layers:

Next is again two fully connected layers with 4096 units. And two dropout layer with rate as 0.5.

```
model.add(Dense(4096, activation='tanh'))
model.add(Dropout(0.5))
model.add(Dense(4096, activation='tanh'))
model.add(Dropout(0.5))
```

dense_1 (Dense)	(None, 4096)	1052672
dropout_1 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 4096)	16781312
dropout_2 (Dropout)	(None, 4096)	0

Output Layer:

Finally, there is a softmax output layer \hat{y} with 200 possible values.

```
model.add(Dense(num_classes, activation='softmax'))
```

dense_3 (Dense)	(None, 200)	819400
-----------------	-------------	--------

Total params: 22,370,760
Trainable params: 22,369,544
Non-trainable params: 1,216

Deploy Model:

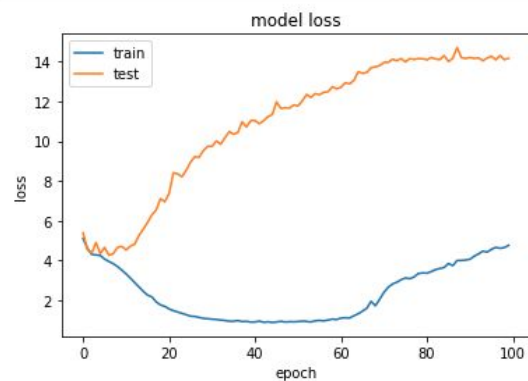

```
hist_1 = model.fit(x_train, y_train,  
                   batch_size=64,  
                   epochs=100,  
                   shuffle=True,  
                   validation_data=(x_test, y_test))
```

Results:

Peak training accuracy at epoch 63 as 0.8134 with 1.1061 training loss;

Test loss: 14.159581217447917



Test accuracy: 0.10366666666666667



Epoch = 50:

Test loss: 11.464195709228516

Test accuracy: 0.1251



Comparing with 100 epochs, the training result looks great! However, it is still over-fitting because the test loss is way higher than the training loss.



Epoch = 500 with only 2 classes data:

Not only shrink the data set to 2 classes but split the validation data from the training data set after finishing an epoch.

```
hist = model.fit(x_train, y_train,
                 batch_size=46,
                 epochs=500,
                 validation_split=0.1,
                 shuffle=True)
```

Test loss: 0.0582424342563

Test accuracy: 0.996



Analysis:

Epoch = 100:

The training accuracy and loss are reasonably tuned before epoch 60 but become weird after that. So we run another training for only 50 epoch below to see if this is because of large number of epoch.

And validation accuracy and loss do not make sense at all. This might because the validation data set is too large which the model does not learn very well.

Epoch = 50:

From the training plot, the convergence is not obvious which means that the epoch number is too low, the curve might be more stable if we run more epoches, like 1000+ epoches. This make perfect sense because our data set is huge and the image complexity is high. And the validation data set is too large comparing with training data set, we should decrease the portion of validation data set and run more epoches. However, because our computer's capacity, we better use AWS or more powerful GPU.

Epoch = 500 with only 2 classes data:

The result reaches our expectation, we get high accuracy and low loss for both train and validation, which proves our assumption that if the number of epoch is high enough, we could get a good result by AlexNet model.

Experiment 4:

Use the pre-trained model as per your team assignment -- InceptionV3

Data:

The same data set as above but the pictures are reshaped with size 150x150 because the minimum of Inception V3 input image's width and length is 75. However, the default input size is 299x299.

Architecture:

Inception V3 model, with weights pre-trained on ImageNet.


```
num_classes = 200

base_model = inception_v3.InceptionV3

base_model = base_model(weights='imagenet', include_top=False)
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1024, activation='relu')(x)
predictions = Dense(num_classes, activation='softmax')(x)
model = Model(inputs=base_model.input, outputs=predictions)
for layer in base_model.layers:
    layer.trainable = False
model.summary()
model.compile(loss='sparse_categorical_crossentropy',
              optimizer=Adam(lr=0.0001),
              metrics=['acc'])
```

Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.5/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5
87916544/87910968 [=====] - 3s 0us/step

Inception (v3)





After the base model of inception v3, we add one globalaveragepooling layer, one 1024 dense layer with 'relu' and a final dense layer of 200, which is the class number, with 'softmax'.

global_average_pooling2d_1 (Glo	(None, 2048)	0	mixed10[0][0]
dense_1 (Dense)	(None, 1024)	2098176	global_average_pooling2d_1[0][0]
dense_2 (Dense)	(None, 200)	205000	dense_1[0][0]
=====			
Total params: 24,105,960			
Trainable params: 2,303,176			
Non-trainable params: 21,802,784			

Result:

Test loss: 16.04263661295573

Test accuracy: 0.004233333333333334



Analysis:

Our training goes well but the validation is not learning at all. Same as part3, the data set is too large and the image is too complex. We need to run more epoches or decrease our data set.

Experiment 5: Transfer Learning

Apply Transfer learning for your team's network to the CIFAR 10 dataset.

Data: Tiny Imagenet 200 dataset as above.


Step 1:

To apply transfer learning to this dataset, it is first of all needed that, to train and save a model with it weights. Accordingly the program uses the CIFAR 10 pre-trained neural network as a base model.

(Please refer to:

https://colab.research.google.com/drive/19OGj3N-_Yateh16qL5seqU6UCaA3F3nN)

It's metrics shows:



```
10000/10000 [=====] - 2s 173us/step
```

```
Test loss: 0.5143789597988129
```

```
Test accuracy: 0.8423
```

Using checkpoints during the training may help the further experiments. The programme saved the model at 50 epochs, which is slightly overfitting.

Step 2:

The programme loaded and cut the model in to required shape, by doing:


```

from keras.models import model_from_json
# load json and create model
json_file = open('model.json', 'r')
loaded_model_json = json_file.read()
json_file.close()
loaded_model = model_from_json(loaded_model_json)
# load weights into new model
loaded_model.load_weights("model.h5")
print("Loaded model from disk")

```

And after doing so, use pop() to get rid of the last layers which makes predictions.
This results in a model like:

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
activation_1 (Activation)	(None, 32, 32, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_2 (Conv2D)	(None, 30, 30, 32)	9248
activation_2 (Activation)	(None, 30, 30, 32)	0
batch_normalization_2 (Batch Normalization)	(None, 30, 30, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 32)	0
spatial_dropout2d_1 (Spatial Dropout)	(None, 15, 15, 32)	0
conv2d_3 (Conv2D)	(None, 15, 15, 64)	18496
activation_3 (Activation)	(None, 15, 15, 64)	0
batch_normalization_3 (Batch Normalization)	(None, 15, 15, 64)	256
conv2d_4 (Conv2D)	(None, 13, 13, 64)	36928
activation_4 (Activation)	(None, 13, 13, 64)	0
batch_normalization_4 (Batch Normalization)	(None, 13, 13, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 64)	0
spatial_dropout2d_2 (Spatial Dropout)	(None, 6, 6, 64)	0
conv2d_5 (Conv2D)	(None, 6, 6, 128)	73856
activation_5 (Activation)	(None, 6, 6, 128)	0
batch_normalization_5 (Batch Normalization)	(None, 6, 6, 128)	512
conv2d_6 (Conv2D)	(None, 4, 4, 128)	147584
activation_6 (Activation)	(None, 4, 4, 128)	0

batch_normalization_6 (Batch Normalization)	(None, 4, 4, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 128)	0
dropout_1 (Dropout)	(None, 2, 2, 128)	0
=====		
Total params: 288,800		
Trainable params: 287,904		
Non-trainable params: 896		

Step 3:

Now that the programme gets a model with certain knowledges about patterns learnt from CIFAR 10 dataset, let us see how these knowledges works for imagenet data.

```
[ ] extracted_features = loaded_model.predict(x_train)
    print(extracted_features.shape)
```

```
↳ (70000, 2, 2, 128)
```

```
▶ validation_features = loaded_model.predict(x_test)
   print(validation_features.shape)
```

```
↳ (30000, 2, 2, 128)
```

By simply walking the imagenet data through the pretrained network, the programme extracts features who are useful for prediction -- CIFAR 10 prediction.

Step 4:

Apply those features into a neural network, and make predictions.

The network is as simple as:


Layer (type)	Output Shape	Param #
=====		
flatten_5 (Flatten)	(None, 3072)	0
dense_9 (Dense)	(None, 512)	1573376
activation_9 (Activation)	(None, 512)	0
dropout_5 (Dropout)	(None, 512)	0

dense_10 (Dense)	(None, 200)	102600
activation_10 (Activation)	(None, 200)	0

=====

Total params: 1,675,976
Trainable params: 1,675,976
Non-trainable params: 0

It takes the input of exactly the same shape as the features have been extracted.
The performance shows:



Last epoch: loss: 3.2950 - acc: 0.2576 - val_loss: 3.9016 - val_acc: 0.1806
The team observed a lack of learning capability in this model. So what might be the problem?

Hypothesis 1: Model. This pretrained model is not able to learn enough features in new data. This can be either because of a low similarity between two datasets, or a low learning capability of model itself.

Hypothesis 2: Data. The data has lost too much information before being applied to the model.


To demonstrate against these hypothesis, this experiment includes the following step.

Step 5:

Hypothesis 1:

The programme introduced a pre-trained vgg-16 model to perform the same steps above. Vgg-16 is trained upon the dataset of imagenet and proved to be a well performed one.

And so gives a result of:




loss: 3.5408 - acc: 0.2347 - val_loss: 5.6596 - val_acc: 0.1260

The learning curve gives a reasonable slope. It is well pre-trained, and thus the slope is steep during the early epochs.

However, since vgg16 is much more complex, and as well proved to be better performed than the first pre-trained model, it is supposed to give a higher accuracy as results. Yet it's metrics tells otherwise. So the team assumed the problem is not mainly about the complexity of the model. It is most likely to be something data-wise.

Hypothesis 2:

The programme increased the training data by a scale of 28.5%. So it gives a result of:



loss: 3.5007 - acc: 0.2275 - val_loss: 3.7130 - val_acc: 0.1950

...which is an increasing in comparison with the first try. The problem can be data-wise. Tiny imagenet gives only 500 images per class, which could hardly be sufficient even for a human brain with 100 billion neurons...

Experiment 6: Fine Tuning

Data: Tiny Imagenet 200 as above.

Step 1:

Load and trim the pre-trained model as it is in Experiment 5.

Step 2:

Freeze the early layers in this model, and train only the late layers. For example:

```
.....  
<keras.layers.core.Activation object at 0x7f4ab492b518> False  
<keras.layers.normalization.BatchNormalization object at 0x7f4ab4888898>  
False  
<keras.layers.convolutional.Conv2D object at 0x7f4ab48ea8d0> True  
<keras.layers.core.Activation object at 0x7f4ab4888908> True  
<keras.layers.normalization.BatchNormalization object at 0x7f4ab47be198>  
True  
<keras.layers.pooling.MaxPooling2D object at 0x7f4ab4807438> True  
<keras.layers.core.Dropout object at 0x7f4ab47be4a8> True
```


Step 3:

Create a model, including the loaded model as its first layers.

```
# Create the model  
model = models.Sequential()  
  
# Add the vgg convolutional base model  
model.add(loaded_model)  
  
# Add new layers  
model.add(layers.Flatten())  
model.add(layers.Dense(1024, activation='relu'))  
model.add(layers.Dropout(0.5))  
model.add(layers.Dense(200, activation='softmax'))
```


During the training, the trainable layers in loaded model it trained along with the additional layers.

The following figures shows the performance when:
5 layers are trainable:



loss: 3.4617 - acc: 0.2250 - val_loss: 3.6515 - val_acc: 0.2071

8 layers are trainable:




15s 164us/step - loss: 3.2330 - acc: 0.2623 - val_loss: 3.4872 - val_acc: 0.2341

13 layers are trainable:

18s 195us/step - loss: 3.0303 - acc: 0.2958 - val_loss: 3.3038 - val_acc: 0.2659

...

And when all layers are released...



```
31s 344us/step - loss: 2.8259 - acc: 0.3322 - val_loss: 2.9612 -  
val_acc: 0.3288
```

An increase in validation scores appears while more layers are released for training, but the training also consumes much more time than ever. There can be a trade-off problem between performance and efficiency when applying fine-tuning.